

COMP_SCI_7318_DLF_A1_Codes

October 4, 2024

1 Import packages

```
[1]: # import packages
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
import matplotlib.pyplot as plt
import seaborn as sns
from torchsummary import summary
import torch
from torch import nn
from torch.utils.data import DataLoader, Dataset
from torchvision import datasets
from torchvision.transforms import ToTensor
from torchvision import models
from torchsummary import summary
import json
from tqdm import tqdm
```

1.0.1 1. Data Preparation

1.1 Load Dataset and Format for Split

```
[2]: # Load dataset from file and format data into a list
with open("diabetes_preprocessed.txt") as diabetes_file:
    rows = diabetes_file.readlines()
row_num = len(rows) # number of observations
col_num = 9         # number of columns
diabetes_array = np.full((row_num, col_num), np.nan)
for row_idx, row in enumerate(rows):
    observation = row.split()
    label = float(observation[0])
    diabetes_array[row_idx, 0] = label
    for feature in observation[1:]:
```

```

        feature_idx, feature_value = feature.split(":")
        diabetes_array[row_idx, int(feature_idx)] = float(feature_value)
diabetes_array[diabetes_array[:, 0] == 1, 0] = 0
diabetes_array[diabetes_array[:, 0] == -1, 0] = 1
diabetes_array = diabetes_array.astype(np.float32)

```

1.2 Split dataset into training and testing

```

[3]: random_seed = 42
bigtrain_set, test_set = train_test_split(diabetes_array, test_size=0.2,
                                         random_state=random_seed,
                                         ↪stratify=diabetes_array[:, 0])
train_set, val_set = train_test_split(bigtrain_set, test_size=0.25,
                                       random_state=random_seed,
                                       ↪stratify=bigtrain_set[:, 0])

```

1.3 Data Cleaning

```

[4]: # Handling missing values
preproc_pl = Pipeline([
    ('imputer', SimpleImputer(strategy='median'))
])

trainp_set = preproc_pl.fit_transform(train_set)
valp_set = preproc_pl.fit_transform(val_set)
testp_set = preproc_pl.fit_transform(test_set)
bigtrainp_set = preproc_pl.fit_transform(bigtrain_set)

# Print dataset details
print('Bigtraining Dataset:', bigtrainp_set.shape)
print('\tTraining Dataset:', trainp_set.shape)
print('\tValidating Dataset:', valp_set.shape)
print('Testing Dataset:', testp_set.shape)

```

```

Bigtraining Dataset: (614, 9)
    Training Dataset: (460, 9)
    Validating Dataset: (154, 9)
Testing Dataset: (154, 9)

```

1.4 Data Overview

```

[5]: def draw_histogram(dataset, data_name, feature_names):
    wid = 4
    hei = int((dataset.shape[1] - 1) / wid)

    fig, axes = plt.subplots(hei, wid,
                             figsize=(int(wid * 5), int(hei * 4)))

    fig.suptitle(f'Histograms of All Features in the {data_name} Set',
    ↪fontsize=16)

```


```

for h in range(hei):
    for w in range(wid):
        k = h * wid + w + 1
        if k >= dataset.shape[1]:
            break
        axes[h, w].hist(dataset[dataset[:, 0] == 0, k], alpha = 0.5, label_
↪ = 'non-diabetes')
        axes[h, w].hist(dataset[dataset[:, 0] == 1, k], alpha = 0.5, label_
↪ = 'diabetes')
        axes[h, w].set_xlabel(feature_names[k - 1])
        axes[h, w].legend(loc='upper right')
    pass


feature_names = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
↪ 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']
draw_histogram(trainp_set, 'Training', feature_names)
draw_histogram(valp_set, 'Validation', feature_names)
draw_histogram(testp_set, 'Testing', feature_names)

```


Histograms of All Features in the Training Set



Histograms of All Features in the Validation Set



Histograms of All Features in the Testing Set



1.5 Create a Custom Dataset

```
[6]: class DiabeteDataset(Dataset):
    def __init__(self, data_array):
        self.diabetes_data = data_array

    def __len__(self):
        return len(self.diabetes_data)

    def __getitem__(self, idx):
        return torch.tensor(self.diabetes_data[idx, 1:]), torch.tensor(self.
        ↪diabetes_data[idx, 0])
```

```

diabetes_train = DiabeteDataset(trainp_set)
diabetes_val = DiabeteDataset(valp_set)
diabetes_test = DiabeteDataset(testp_set)
diabetes_bigtrain = DiabeteDataset(bigtrainp_set)

```

1.0.2 2. Training Models: Multi-scale Random Search

2.1 Baseline

```

[7]: class MyScaleLayer(nn.Module):
    def __init__(self):
        super().__init__()

    def forward(self, input):
        return (input + 1.0) * 0.5

class DNNRunner():
    def __init__(self, train_dataset=None, val_dataset=None, hyper_params=None):
        if hyper_params is None:
            hyper_params = {}
        if 'model_version' not in hyper_params:
            hyper_params['model_version'] = 'single'
        if 'actfn' not in hyper_params:
            hyper_params['actfn'] = 'relu'
        if 'batch_size' not in hyper_params:
            hyper_params['batch_size'] = 16
        if 'learning_rate' not in hyper_params:
            hyper_params['learning_rate'] = 1.0
        if 'optimizer' not in hyper_params:
            hyper_params['optimizer'] = 'sgd'

        self.trial_tag = '_'.join(map(str, [
            hyper_params['model_version'],
            hyper_params['actfn'],
            hyper_params['batch_size'],
            hyper_params['learning_rate'],
            hyper_params['optimizer']
        ]))

        self.device = self.get_device()
        self.batch_size = int(hyper_params['batch_size'])
        self.learning_rate = hyper_params['learning_rate']
        self.actfn = self.get_actfn(hyper_params['actfn'])
        self.model_version = hyper_params['model_version']

```

```

        if train_dataset is not None:
            self.train_dataloader = DataLoader(train_dataset, batch_size=self.
↪batch_size, shuffle=True)
        if val_dataset is not None:
            self.val_dataloader = DataLoader(val_dataset, batch_size=self.
↪batch_size, shuffle=False)

        self.model = self.create_model().to(self.device)
        self.loss_fn = nn.BCELoss() # nn.BCEWithLogitsLoss()
        self.optimizer = self.get_optimizer(hyper_params['optimizer'])(self.
↪model.parameters(), lr=self.learning_rate)

    @staticmethod
    def get_device():
        device = (
            "cuda"
            if torch.cuda.is_available()
            else "mps"
            if torch.backends.mps.is_available()
            else "cpu"
        )
        return device

    @staticmethod
    def get_optimizer(optimizer_name):
        if optimizer_name == 'sgd':
            return torch.optim.SGD
        elif optimizer_name == 'adam':
            return torch.optim.Adam
        elif optimizer_name == 'rmsprop':
            return torch.optim.RMSprop
        elif optimizer_name == 'nadam':
            return torch.optim.NAdam
        else:
            raise NotImplementedError(f'Unknown optimizer {optimizer_name}')

    @staticmethod
    def get_actfn(actfn):
        if actfn == 'relu':
            return nn.ReLU
        elif actfn == 'elu':
            return nn.ELU
        elif actfn == 'selu':
            return nn.SELU
        elif actfn == 'leaky_relu':
            return nn.LeakyReLU
        else:

```

```

        raise NotImplementedError(f'Unknown actfn {actfn}')

    @classmethod
    def report_model(cls, model_version):
        print(f'{model_version}:')
        single_model = cls(hyper_params={'model_version': model_version}).
        ↪ create_model()
        summary(single_model, (1, 8))

    def create_model(self):

        if self.model_version == 'single':
            model = nn.Sequential(
                nn.Linear(8, 1),
                nn.Tanh(),
                MyScaleLayer()
            )

        elif self.model_version == 'shallow':
            model = nn.Sequential(
                nn.Linear(8, 16),
                self.actfn(),
                nn.Linear(16, 16),
                self.actfn(),
                nn.Linear(16, 16),
                self.actfn(),
                nn.Linear(16, 16),
                self.actfn(),
                nn.Linear(16, 1),
                nn.Tanh(),
                MyScaleLayer()
            )

        elif self.model_version == 'deep':
            model = nn.Sequential(
                nn.Linear(8, 16),
                self.actfn(),
                nn.Linear(16, 16),
                self.actfn(),
                nn.Linear(16, 16),
                self.actfn(),
                nn.Linear(16, 32),
                self.actfn(),
                nn.Linear(32, 32),
                self.actfn(),
                nn.Linear(32, 16),
                self.actfn(),
                nn.Linear(16, 16),

```

```

        self.actfn(),
        nn.Linear(16, 16),
        self.actfn(),
        nn.Linear(16, 1),
        nn.Tanh(),
        MyScaleLayer()
    )
    else:
        raise NotImplementedError(f'Unknown model version {self.
↪model_version}.'.)

    return model.to(self.device)

def run(self, epoch):
    result = {
        'trial_tag': self.trial_tag,
        'epochs': epoch,
        'train_losses': [],
        'train_accs': [],
        'val_accs': [],
    }
    with tqdm(total=epoch) as pbar:
        for _ in range(epoch):
            pbar.update(1)
            self.model.train()
            train_loss, train_acc = self.train_one_epoch()
            result['train_losses'].append(train_loss)
            result['train_accs'].append(train_acc)

            self.model.eval()
            val_acc, all_y, all_y_pred = self.validate_one_epoch()
            result['val_accs'].append(val_acc)
            pbar.set_description(f'[t_loss]{train_loss:.2f};[v_acc]{val_acc:
↪.2%}')
```

Evaluate the last

```

metrics = self.compute_metrics(all_y, all_y_pred)
return **metrics, **result}

def train_one_epoch(self):
    train_loss = 0.0
    train_acc = 0.0
    for X, y in self.train_dataloader:
        X, y = X.to(self.device), y.to(self.device)
        pred = self.model(X).reshape(-1)
        loss = self.loss_fn(pred, y)
        loss.backward()

```



```

        self.optimizer.step()
        self.optimizer.zero_grad()
        train_loss += loss.item()
        pred_y = (pred.detach().cpu() >= 0.5).int()
        train_acc += accuracy_score(pred_y, y.detach().cpu())
    train_loss /= len(self.train_dataloader)
    train_acc /= len(self.train_dataloader)
    return train_loss, train_acc

def validate_one_epoch(self, test_dataset=None):
    test_dataloader = (
        DataLoader(test_dataset, batch_size=self.batch_size, shuffle=False)
        if test_dataset is not None
        else self.val_dataloader
    )
    all_y, all_pred = [], []
    with torch.no_grad():
        for X, y in test_dataloader:
            X, y = X.to(self.device), y.to(self.device)
            pred = self.model(X).reshape(-1)
            all_y.append(y.detach().cpu())
            all_pred.append(pred.detach().cpu())
    all_y, all_pred = torch.cat(all_y), torch.cat(all_pred)
    all_y_pred = (all_pred >= 0.5).int()
    return accuracy_score(all_y, all_y_pred), all_y, all_y_pred

def compute_metrics(self, all_y, all_y_pred):
    acc = accuracy_score(all_y, all_y_pred)
    cr = classification_report(all_y, all_y_pred,
                              output_dict=True, zero_division=1)
    avg_results = cr.get('weighted avg', {})
    cm = confusion_matrix(all_y, all_y_pred)
    metrics = {
        'val_acc': acc,
        'precision': avg_results.get('precision', None),
        'recall': avg_results.get('recall', None),
        'f1_score': avg_results.get('f1-score', None),
        'confusion_matrix': json.dumps(cm.tolist(), default=lambda x:
↪float(x)),
    }
    return metrics

```

```

[9]: DNNRunner.report_model('single')
DNNRunner.report_model('shallow')
DNNRunner.report_model('deep')

```

single:

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1, 1]	9
Tanh-2	[-1, 1, 1]	0
MyScaleLayer-3	[-1, 1, 1]	0

Total params: 9

Trainable params: 9

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.00

Params size (MB): 0.00

Estimated Total Size (MB): 0.00

shallow:

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1, 16]	144
ReLU-2	[-1, 1, 16]	0
Linear-3	[-1, 1, 16]	272
ReLU-4	[-1, 1, 16]	0
Linear-5	[-1, 1, 16]	272
ReLU-6	[-1, 1, 16]	0
Linear-7	[-1, 1, 16]	272
ReLU-8	[-1, 1, 16]	0
Linear-9	[-1, 1, 1]	17
Tanh-10	[-1, 1, 1]	0
MyScaleLayer-11	[-1, 1, 1]	0

Total params: 977

Trainable params: 977

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.00

Params size (MB): 0.00

Estimated Total Size (MB): 0.00

deep:

Layer (type)	Output Shape	Param #
Linear-1	[-1, 1, 16]	144
ReLU-2	[-1, 1, 16]	0
Linear-3	[-1, 1, 16]	272
ReLU-4	[-1, 1, 16]	0

Linear-5	[-1, 1, 16]	272
ReLU-6	[-1, 1, 16]	0
Linear-7	[-1, 1, 32]	544
ReLU-8	[-1, 1, 32]	0
Linear-9	[-1, 1, 32]	1,056
ReLU-10	[-1, 1, 32]	0
Linear-11	[-1, 1, 16]	528
ReLU-12	[-1, 1, 16]	0
Linear-13	[-1, 1, 16]	272
ReLU-14	[-1, 1, 16]	0
Linear-15	[-1, 1, 16]	272
ReLU-16	[-1, 1, 16]	0
Linear-17	[-1, 1, 1]	17
Tanh-18	[-1, 1, 1]	0
MyScaleLayer-19	[-1, 1, 1]	0

Total params: 3,377
Trainable params: 3,377
Non-trainable params: 0

Input size (MB): 0.00
Forward/backward pass size (MB): 0.00
Params size (MB): 0.01
Estimated Total Size (MB): 0.02

```
[10]: baseline_params = {
    'model_version': 'single',
    'actfn': 'relu',
    'batch_size': 16,
    'learning_rate': 1e-4,
    'optimizer': 'adam'
}

dnn_baseline = DNNRunner(diabetes_bigtrain, diabetes_test, baseline_params)
baseline_res = dnn_baseline.run(200)
```

```
[t_loss]22.40;[v_acc]53.90%: 100%|      | 200/200 [00:17<00:00, 11.14it/s]
```

2.2 Multi-scale random search for hyper parameters We do 50 training runs in total to balance the search area and the efficiency.

- 30 training runs in the first random searching
- 20 training runs in the second searching

```
[11]: def random_search(hyper_params, trial_num=30):

    all_res = []
```

```

# Randomly select hyperparameters
used_config = set()
all_params = []
for trial_idx in range(trial_num):
    dict_params = {}
    while True:
        config_record = []
        for key in hyper_params.keys():
            dict_params[key] = np.random.choice(hyper_params[key])
            config_record.append(dict_params[key])
        config_record = tuple(config_record)
        if config_record not in used_config:
            used_config.add(config_record)
            break
    all_params.append(dict_params)

# Run trials
for idx, params in enumerate(all_params):
    dnn_trial = DNNRunner(diabetes_train, diabetes_val, params)
    print(f'Trail {idx} ({dnn_trial.trial_tag}):')
    res = dnn_trial.run(200)
    # print(f'Trail {idx} completed.')
    all_res.append(res)

return all_res

```

```

[14]: # First search
first_search_params = {
    'model_version': ['single', 'shallow', 'deep'],
    'actfn': ['relu', 'elu', 'selu', 'leaky_relu'],
    'batch_size': [8, 16, 32],
    'learning_rate': [0.1, 0.01, 0.001, 0.0001],
    'optimizer': ['adam', 'rmsprop', 'nadam']
} # 3 * 4 * 3 * 4 * 3 = 432

first_results = random_search(first_search_params)

```

Trail 0 (deep_elu_16_0.01_nadam):

[t_loss]0.42;[v_acc]75.32%: 100% | 200/200 [00:29<00:00, 6.80it/s]

Trail 1 (shallow_elu_8_0.0001_nadam):

[t_loss]0.46;[v_acc]79.22%: 100% | 200/200 [00:39<00:00, 5.11it/s]

Trail 2 (single_selu_32_0.01_rmsprop):

[t_loss]34.38;[v_acc]64.94%: 100% | 200/200 [00:07<00:00, 25.46it/s]

Trail 3 (shallow_selu_8_0.001_adam):

[t_loss]0.32;[v_acc]74.68%: 100%| | 200/200 [00:36<00:00, 5.46it/s]

Trail 4 (single_relu_16_0.0001_nadam):

[t_loss]65.05;[v_acc]35.06%: 100%| | 200/200 [00:13<00:00, 14.96it/s]

Trail 5 (shallow_elu_16_0.0001_adam):

[t_loss]0.48;[v_acc]76.62%: 100%| | 200/200 [00:19<00:00, 10.23it/s]

Trail 6 (deep_relu_8_0.01_adam):

[t_loss]0.65;[v_acc]64.94%: 100%| | 200/200 [00:50<00:00, 3.93it/s]

Trail 7 (deep_relu_16_0.001_nadam):

[t_loss]0.30;[v_acc]74.68%: 100%| | 200/200 [00:27<00:00, 7.30it/s]

Trail 8 (deep_selu_16_0.01_nadam):

[t_loss]0.65;[v_acc]64.94%: 100%| | 200/200 [00:27<00:00, 7.34it/s]

Trail 9 (single_leaky_relu_32_0.0001_nadam):

[t_loss]23.21;[v_acc]63.64%: 100%| | 200/200 [00:08<00:00, 23.76it/s]

Trail 10 (deep_leaky_relu_32_0.001_adam):

[t_loss]0.31;[v_acc]75.32%: 100%| | 200/200 [00:14<00:00, 13.63it/s]

Trail 11 (shallow_selu_16_0.01_nadam):

[t_loss]0.42;[v_acc]74.68%: 100%| | 200/200 [00:20<00:00, 9.71it/s]

Trail 12 (single_leaky_relu_32_0.01_adam):

[t_loss]33.55;[v_acc]64.94%: 100%| | 200/200 [00:07<00:00, 25.32it/s]

Trail 13 (deep_leaky_relu_8_0.01_nadam):

[t_loss]64.87;[v_acc]35.06%: 100%| | 200/200 [00:52<00:00, 3.82it/s]

Trail 14 (deep_relu_8_0.1_rmsprop):

[t_loss]34.91;[v_acc]64.94%: 100%| | 200/200 [00:49<00:00, 4.05it/s]

Trail 15 (deep_leaky_relu_8_0.001_rmsprop):

[t_loss]0.22;[v_acc]77.27%: 100%| | 200/200 [00:48<00:00, 4.09it/s]

Trail 16 (deep_selu_8_0.01_rmsprop):

[t_loss]64.87;[v_acc]35.06%: 100%| | 200/200 [00:49<00:00, 4.06it/s]

Trail 17 (deep_leaky_relu_32_0.0001_rmsprop):

[t_loss]0.51;[v_acc]76.62%: 100%| | 200/200 [00:14<00:00, 13.41it/s]

Trail 18 (single_selu_16_0.01_nadam):

[t_loss]34.84;[v_acc]64.94%: 100%| | 200/200 [00:14<00:00, 13.87it/s]

Trail 19 (shallow_elu_16_0.01_adam):

```

[t_loss]0.45;[v_acc]78.57%: 100%|      | 200/200 [00:20<00:00,  9.74it/s]
Trail 20 (deep_relu_32_0.0001_adam):
[t_loss]0.54;[v_acc]76.62%: 100%|      | 200/200 [00:15<00:00, 12.63it/s]
Trail 21 (deep_relu_32_0.1_rmsprop):
[t_loss]34.72;[v_acc]64.94%: 100%|      | 200/200 [00:15<00:00, 13.23it/s]
Trail 22 (shallow_leaky_relu_8_0.0001_adam):
[t_loss]0.50;[v_acc]74.03%: 100%|      | 200/200 [00:38<00:00,  5.25it/s]
Trail 23 (shallow_elu_32_0.1_rmsprop):
[t_loss]34.72;[v_acc]64.94%: 100%|      | 200/200 [00:11<00:00, 17.17it/s]
Trail 24 (shallow_leaky_relu_16_0.0001_rmsprop):
[t_loss]0.54;[v_acc]72.08%: 100%|      | 200/200 [00:19<00:00, 10.09it/s]
Trail 25 (deep_selu_8_0.0001_rmsprop):
[t_loss]0.41;[v_acc]79.87%: 100%|      | 200/200 [00:49<00:00,  4.05it/s]
Trail 26 (shallow_leaky_relu_8_0.001_nadam):
[t_loss]0.31;[v_acc]75.32%: 100%|      | 200/200 [00:42<00:00,  4.66it/s]
Trail 27 (single_leaky_relu_8_0.001_rmsprop):
[t_loss]0.74;[v_acc]64.29%: 100%|      | 200/200 [00:24<00:00,  8.13it/s]
Trail 28 (single_leaky_relu_8_0.1_adam):
[t_loss]65.30;[v_acc]35.06%: 100%|      | 200/200 [00:25<00:00,  7.97it/s]
Trail 29 (shallow_elu_8_0.001_adam):
[t_loss]0.33;[v_acc]75.97%: 100%|      | 200/200 [00:38<00:00,  5.13it/s]

```

```

[15]: first_result_df = pd.DataFrame(first_results)
      first_result_sorted = first_result_df.sort_values('val_acc', ascending=False)
      first_result_sorted.head(n=3)

```

```

[15]:   val_acc  precision    recall  f1_score   confusion_matrix \
25  0.798701   0.798202   0.798701   0.789490  [[92, 8], [23, 31]]
1   0.792208   0.788371   0.792208   0.786554  [[89, 11], [21, 33]]
19  0.785714   0.785434   0.785714   0.774092  [[92, 8], [25, 29]]

```

```

          trial_tag  epochs \
25  deep_selu_8_0.0001_rmsprop    200
1   shallow_elu_8_0.0001_nadam    200
19   shallow_elu_16_0.01_adam    200

```

```

          train_losses \

```

```

25 [0.6738769997810495, 0.6371511539508556, 0.618...
1  [0.7707187271323698, 0.7284242238464027, 0.698...
19 [0.6834645076044674, 0.6346456274904054, 0.607...

```

```

                                train_accs \
25 [0.5991379310344828, 0.6594827586206896, 0.663...
1  [0.5969827586206896, 0.6379310344827587, 0.648...
19 [0.6185344827586207, 0.6580459770114943, 0.687...

```

```

                                val_accs
25 [0.6493506493506493, 0.6753246753246753, 0.694...
1  [0.6688311688311688, 0.6623376623376623, 0.668...
19 [0.6493506493506493, 0.7662337662337663, 0.714...

```

[16]: *# Second Search*

```

second_hyper_params = {
    'model_version': ['shallow', 'deep'],
    'actfn': ['selu', 'elu'],
    'batch_size': [8, 16],
    'learning_rate': [0.01, 0.0001],
    'optimizer': ['adam', 'nadam', 'rmsprop']
} # 2 * 2 * 2 * 2 * 3 = 64

second_results = random_search(second_hyper_params, trial_num=20)

```

Trail 0 (deep_selu_8_0.0001_rmsprop):

```
[t_loss]0.41;[v_acc]76.62%: 100%|      | 200/200 [00:50<00:00, 3.98it/s]
```

Trail 1 (deep_selu_8_0.0001_nadam):

```
[t_loss]0.41;[v_acc]76.62%: 100%|      | 200/200 [00:53<00:00, 3.75it/s]
```

Trail 2 (deep_selu_16_0.01_adam):

```
[t_loss]0.47;[v_acc]77.27%: 100%|      | 200/200 [00:26<00:00, 7.48it/s]
```

Trail 3 (deep_elu_8_0.01_nadam):

```
[t_loss]0.43;[v_acc]80.52%: 100%|      | 200/200 [00:53<00:00, 3.76it/s]
```

Trail 4 (shallow_elu_16_0.01_adam):

```
[t_loss]0.47;[v_acc]76.62%: 100%|      | 200/200 [00:20<00:00, 9.89it/s]
```

Trail 5 (deep_elu_16_0.01_nadam):

```
[t_loss]0.44;[v_acc]80.52%: 100%|      | 200/200 [00:27<00:00, 7.16it/s]
```

Trail 6 (deep_selu_16_0.01_rmsprop):

```
[t_loss]0.46;[v_acc]78.57%: 100%|      | 200/200 [00:25<00:00, 7.73it/s]
```

Trail 7 (shallow_selu_16_0.0001_adam):

```

[t_loss]0.45;[v_acc]77.27%: 100%|      | 200/200 [00:20<00:00, 9.90it/s]
Trail 8 (shallow_selu_16_0.0001_rmsprop):
[t_loss]0.46;[v_acc]79.22%: 100%|      | 200/200 [00:20<00:00, 9.54it/s]
Trail 9 (deep_selu_8_0.01_adam):
[t_loss]0.48;[v_acc]79.22%: 100%|      | 200/200 [00:51<00:00, 3.92it/s]
Trail 10 (shallow_selu_16_0.0001_nadam):
[t_loss]0.45;[v_acc]81.17%: 100%|      | 200/200 [00:21<00:00, 9.47it/s]
Trail 11 (shallow_selu_8_0.01_rmsprop):
[t_loss]0.48;[v_acc]67.53%: 100%|      | 200/200 [00:36<00:00, 5.43it/s]
Trail 12 (deep_selu_16_0.0001_rmsprop):
[t_loss]0.43;[v_acc]77.27%: 100%|      | 200/200 [00:25<00:00, 7.70it/s]
Trail 13 (shallow_elu_8_0.01_rmsprop):
[t_loss]0.47;[v_acc]76.62%: 100%|      | 200/200 [00:36<00:00, 5.42it/s]
Trail 14 (deep_selu_8_0.0001_adam):
[t_loss]0.42;[v_acc]79.87%: 100%|      | 200/200 [00:50<00:00, 3.97it/s]
Trail 15 (deep_elu_16_0.0001_rmsprop):
[t_loss]0.45;[v_acc]75.97%: 100%|      | 200/200 [00:26<00:00, 7.65it/s]
Trail 16 (shallow_selu_16_0.01_adam):
[t_loss]0.45;[v_acc]74.68%: 100%|      | 200/200 [00:20<00:00, 9.76it/s]
Trail 17 (shallow_elu_8_0.0001_rmsprop):
[t_loss]0.44;[v_acc]81.17%: 100%|      | 200/200 [00:37<00:00, 5.39it/s]
Trail 18 (deep_elu_16_0.01_adam):
[t_loss]0.46;[v_acc]80.52%: 100%|      | 200/200 [00:26<00:00, 7.43it/s]
Trail 19 (deep_selu_8_0.01_nadam):
[t_loss]0.65;[v_acc]64.94%: 100%|      | 200/200 [00:53<00:00, 3.73it/s]

```

```

[17]: second_result_df = pd.DataFrame(second_results)
      second_result_sorted = second_result_df.sort_values('val_acc', ascending=False)
      second_result_sorted.head(n=3)

```

```

[17]:   val_acc  precision  recall  f1_score  confusion_matrix \
17  0.811688   0.808888  0.811688  0.807183  [[90, 10], [19, 35]]
10  0.811688   0.809608  0.811688  0.805919  [[91, 9], [20, 34]]
3   0.805195   0.809487  0.805195  0.793757  [[94, 6], [24, 30]]

```



```

trial_tag epochs \
17 shallow_elu_8_0.0001_rmsprop      200
10 shallow_selu_16_0.0001_nadam      200
3    deep_elu_8_0.01_nadam           200

train_losses \
17 [0.7372722641147417, 0.6481008534801418, 0.634...
10 [0.693741127334792, 0.6478117252218312, 0.6247...
3  [0.6579945817075926, 0.6532618002644901, 0.651...

train_accs \
17 [0.5474137931034483, 0.6357758620689655, 0.640...
10 [0.646551724137931, 0.6185344827586207, 0.6652...
3  [0.6379310344827587, 0.6422413793103449, 0.642...

val_accs
17 [0.6168831168831169, 0.6168831168831169, 0.642...
10 [0.6298701298701299, 0.6623376623376623, 0.649...
3  [0.6493506493506493, 0.6818181818181818, 0.649...

```

2.3 Best Model and Comparison with Baseline

[18]: *# According the output above, the best model has following hyper parameters:*

```

best_hyper_params = {
    'model_version': 'shallow',
    'actfn': 'elu',
    'batch_size': 8,
    'learning_rate': 0.0001,
    'optimizer': 'rmsprop'
}

dnn_best = DNNRunner(diabetes_bigtrain, diabetes_test, best_hyper_params)
best_param_res = dnn_best.run(200)

```

```
[t_loss]0.43;[v_acc]71.43%: 100%|      | 200/200 [00:48<00:00, 4.13it/s]
```

Comparison between baseline and best model

[19]: `comp_base_best_df = pd.DataFrame([baseline_res, best_param_res])`
`comp_base_best_df`

```

[19]:   val_acc  precision    recall  f1_score   confusion_matrix \
0  0.538961  0.480404  0.538961  0.500361  [[76, 24], [47, 7]]
1  0.714286  0.702878  0.714286  0.700023  [[86, 14], [30, 24]]

trial_tag epochs \
0    single_relu_16_0.0001_adam      200
1  shallow_elu_8_0.0001_rmsprop      200

```

```

                                train_losses \
0  [47.500784213726334, 47.656470323220276, 47.33...
1  [0.7423065556334211, 0.6635746069542774, 0.633...

                                train_accs \
0  [0.40170940170940167, 0.4081196581196581, 0.41...
1  [0.5573593073593073, 0.6282467532467533, 0.668...

                                val_accs
0  [0.461038961038961, 0.461038961038961, 0.46103...
1  [0.6103896103896104, 0.7012987012987013, 0.675...

```

1.0.3 3. Result visualization

3.1 Learning curve

```

[20]: #
      # Drawing the learning curves
      #
      def drawing_curves(res_list):
          wid = 5
          hei = int(np.ceil(len(res_list) / wid))
          fig, axes = plt.subplots(hei, wid,
                                   figsize=(int(wid * 5), int(hei * 4)),
                                   sharex=True, sharey=False,
                                   squeeze=False)

          for h in range(hei):
              for w in range(wid):
                  k = h * wid + w
                  if k >= len(res_list):
                      break
                  loss_train = res_list[k]['train_losses']
                  loss_valid = res_list[k]['val_accs']
                  epochs = list(range(1, res_list[k]['epochs'] + 1))
                  axes[h, w].set_title(res_list[k]['trial_tag'])
                  axes[h, w].plot(epochs, loss_train, label='Train')
                  axes[h, w].plot(epochs, loss_valid, label='Validation')
                  axes[h, w].set_xlabel('Epoch nums')
                  axes[h, w].set_ylabel('Loss')
                  axes[h, w].legend(loc='upper right')

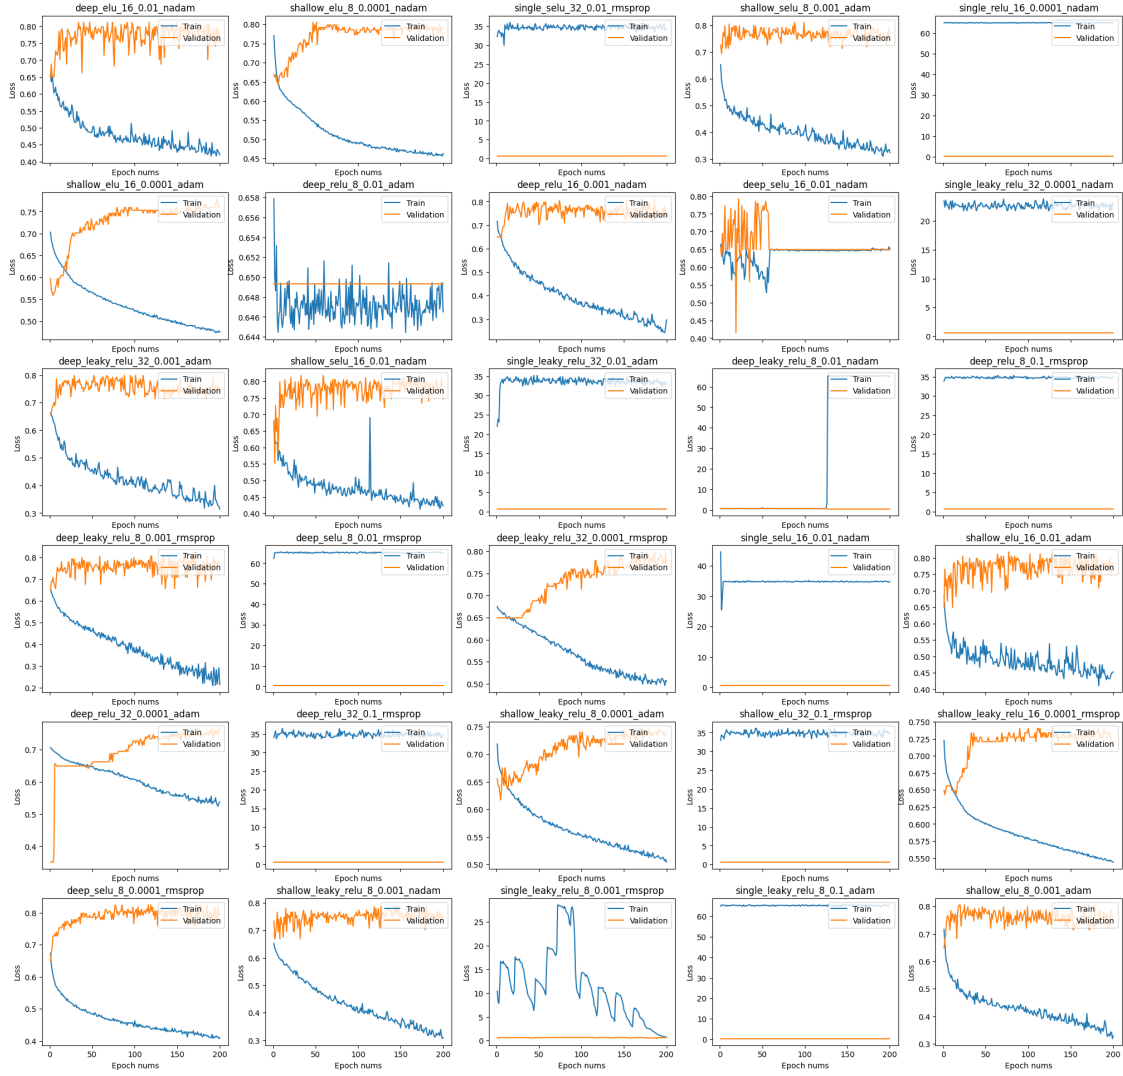
          pass

```

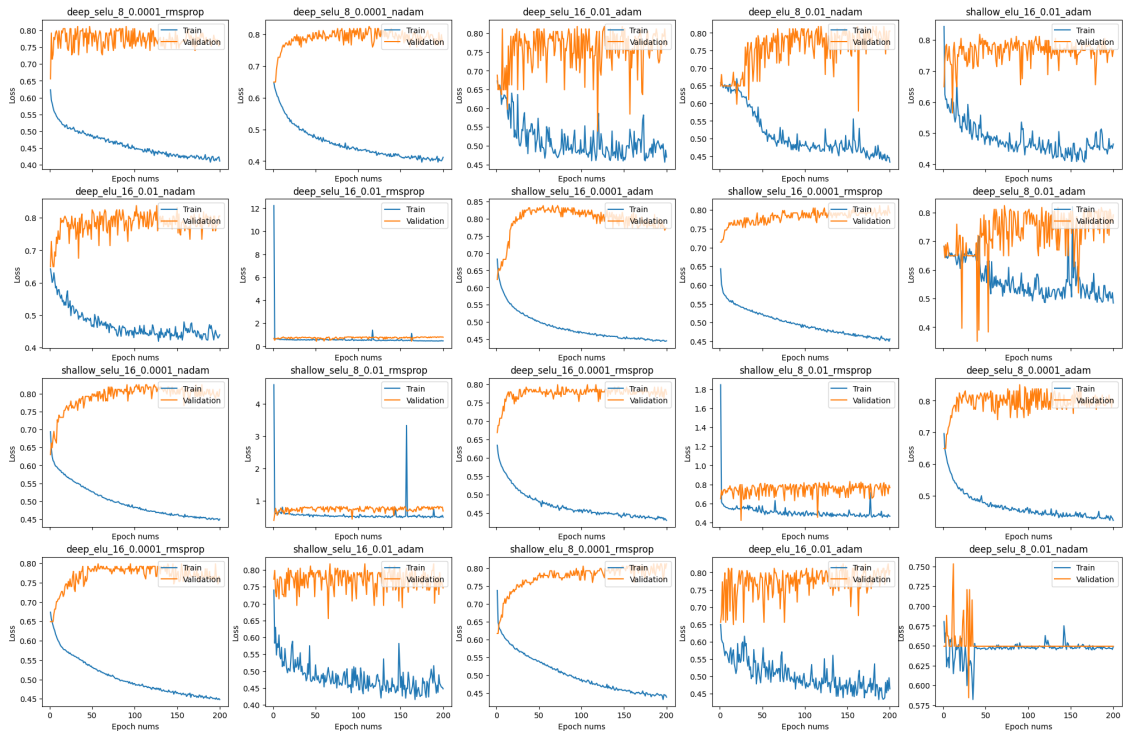
```

[21]: drawing_curves(first_results)

```



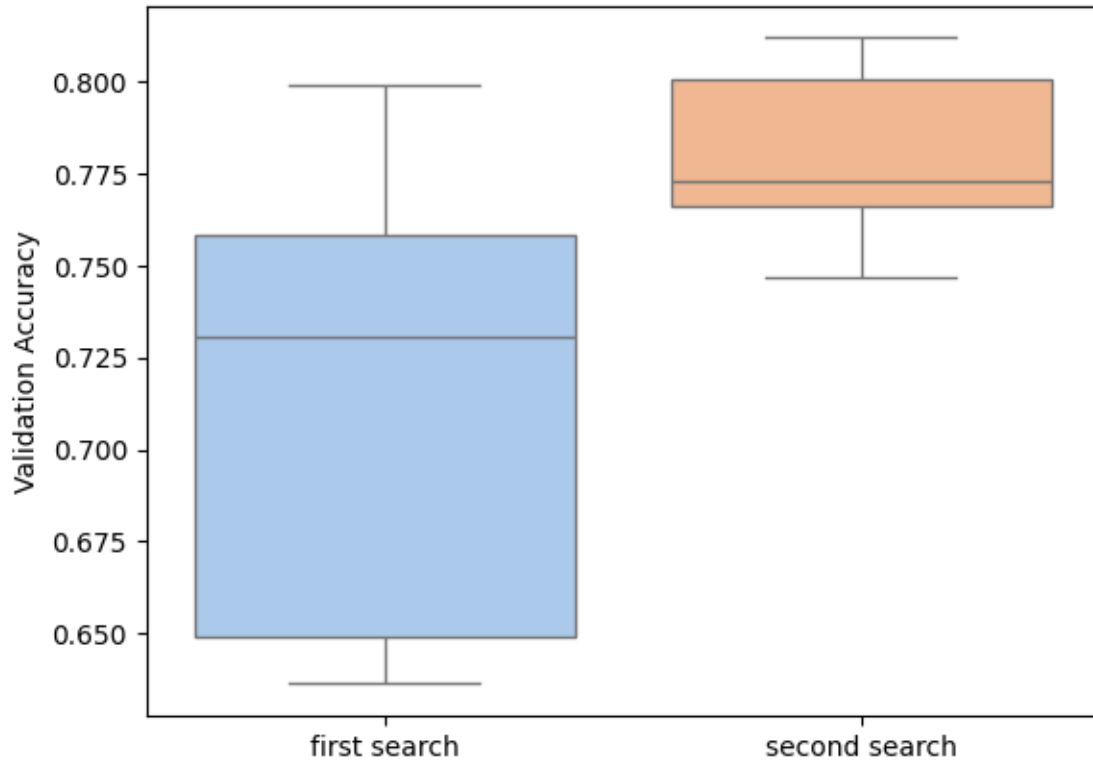
[22] : `drawing_curves(second_results)`



3.2 boxplots

```
[23]: # boxplots to show
#
# Drawing the learning curves
#
first_result_df['tag'] = "first search"
second_result_df['tag'] = "second search"
all_result_df = pd.concat([first_result_df, second_result_df])


ax = sns.boxplot(data=all_result_df, x='tag', hue="tag", y="val_acc",
                  orient='v', showfliers=False, palette=sns.
                  color_palette('pastel', 2))
ax.set(xlabel=None, ylabel='Validation Accuracy')
pass
```



3.3 Training Acc v.s. Validating Acc

```
[24]: def draw_accs(res_df, title=None):
    fig = plt.figure(figsize=(4, 4))
    # res_df['train_acc'] = res_df['trains_accs']
    res_df['train_acc'] = res_df['train_accs'].apply(lambda x: x[-1])
    res_df['model_version'] = res_df['trial_tag'].apply(lambda x: x.
    ↪split('_')[0])
    plt.plot((0, 1), (0, 1), ls='--', c='k')
    sns.scatterplot(data=res_df, x='train_acc', y='val_acc',
    ↪hue='model_version')
    if title is not None:
        plt.title(title)
    plt.xlabel('Training Accuracy')
    plt.ylabel('Validation Accuracy')
    pass

[25]: draw_accs(pd.concat([first_result_df, second_result_df]), 'Accuracy
    ↪distribution')
    # draw_accs(second_result_df, 'Accuracy distribution in second search')
```



3.4 Confusion matrix

```
[26]: def draw_confusion(conf_mat, title=None):
    conf_mat = json.loads(conf_mat)
    plt.figure(figsize=(5, 3))
    if title is not None:
        plt.title(title)
    sns.heatmap(conf_mat, annot=True, fmt='d')
    plt.xlabel('Pred')
    plt.ylabel('Truth')
    plt.show()

draw_confusion(baseline_res['confusion_matrix'], "Confusion matrix of the_
↳Baseline Model")
draw_confusion(best_param_res['confusion_matrix'], "Confusion matrix of the_
↳Improved Model")
```

